


Workstream#2 Data & Algorithm Updates (Sep-09-2025)

LF InfiniEdge AI Workstream#2: data and algorithms

- Focus on AI applications edge computing & use cases (technological & business)
- Pushing the boundary of LLMs with data (on-prem) and algorithms
 - Retrieval accuracy 20% ⇒ 70+% (Complex RAG KDD 2024)
 - Coding capability 25% ⇒ 45% (Reflections) ⇒ 90% (Human Hints) (AI Competitive Coding Neurips 2024)
- Edge Data Agent – Data-on-prem, Code-on-the-Fly



Previous Key Learnings – SFT, RAG & Reasoning

Foundational model dominates the performance (trust scaling law)

SFT change styles and format but not adding new knowledges

RAG (2024 KDD Cup Complex Retrieval Challenges)

- Accuracy 20% \Rightarrow 70+% (Fine Tune / Data Preproc.)
- Battling with the model's memory

Reasoning (2024 NeuRIPs AI for Competitive Coding)

- Accuracy 25% \Rightarrow 45% (Reflections) \Rightarrow 90% (Human Hints)
- Agents x50 cost

EDA converts a user's own **data into** a *local*,
on-demand **agent** service at the **edge**

Data-on-prem, Code-on-the-Fly



Leaderboards

SWE-agent-LM-32B, the open-weight SotA on Verified, trained on synthetic data generated by **SWE-smith**. [More in the paper!](#)

Bash Only Verified Lite Full Multimodal


Bash Only evaluates all LMs with a [minimal agent](#) on SWE-bench Verified ([details](#))

Filters: [Open Scaffold ▾](#) [All Tags ▾](#)

Model	% Resolved	Org	Date	Logs	Trajs	Site	Release
 Claude 4 Opus (20250514)	67.60		2025-08-02	✓	✓	🔗	-
 GPT-5 (2025-08-07) (medium reasoning)	65.00		2025-08-07	✓	✓	🔗	-
 Claude 4 Sonnet (20250514)	64.93		2025-07-26	✓	✓	🔗	-
 GPT-5 mini (2025-08-07) (medium reasoning)	59.80		2025-08-07	✓	✓	🔗	-
 o3 (2025-04-16)	58.40		2025-07-26	✓	✓	🔗	-
 Qwen3-Coder 480B/A35B Instruct	55.40		2025-08-02	✓	✓	🔗	-
 Gemini 2.5 Pro (2025-05-06)	53.60		2025-07-26	✓	✓	🔗	-
 Claude 3.7 Sonnet (20250219)	52.80		2025-07-20	✓	✓	🔗	-
 o4-mini (2025-04-16)	45.00		2025-07-26	✓	✓	🔗	-
 Kimi K2 Instruct	43.80		2025-08-07	✓	✓	🔗	-

(Previously) Single Agent.yaml

```
src > prompts > ! custom_agent.yaml
1   system_prompt: |-
2     You are an expert data analysis assistant who can perform any type of data exploration, processing, or transformation task using co
3     You can leverage Python code within your responses to operate on data, explore patterns, produce visualizations, and answer complex
4
5     You have access to a set of specialized tools (Python functions) which you can call by providing a code block in your response.
6     The steps to solve any problem must follow these sequences:
7
8     1. **Thought**:
9       | Provide a concise explanation of your reasoning and the approach or tools you're about to use.
10
11    2. **Code**:
12      | Provide a Python code snippet (enclosed in triple backticks) ending with `<end_code>`.
13      - Within these code snippets, you can call any relevant tool or write standard Python code.
14      - The output of your code snippet is then displayed in the **"Observation"** field of the next step.
15
16    3. **Observation**:
17      | Show and summarize the results from your code snippet execution.
18
19      You can do multiple cycles of Thought → Code → Observation if needed, until you have a definitive solution.
20
21    **Final Answer**:
22      Once you have enough information to produce your conclusion or final data insights, provide your answer using the `final_answer` tool
23
24      ---
25
26    **Example**:
27      **Task**: "Write SQL code to retrieve information from the SQLite database `data/mydata.db`."
28
29      - **Thought**:
30        | "I will import the `SQLiteTool`, run the agent with the query, and then provide the result."
31
32      - **Code**:
33        ````py
34          from smolagents import ToolCallingAgent, LiteLLMModel
35          from tools.tool_sqlite import SQLiteTool
```



Initial plan is to follow the smolagent template to craft a single (all-in-one) system prompt for edge data agent and together with hardcoded functions and tools (e.g. data connectors, mcp server, ...)

Git Update: adding EDA Eval Suite

Evaluation Results

Task	General LLM	Local Agent	Δ (Agent – LLM)
p1_finance_invoice_match	0.33	1.00	0.67
p1_hr_post_termination	0.00	1.00	1.00
p1_ops_spike	0.00	1.00	1.00
p2_emails_discount_thread	0.00	1.00	1.00
p2_audio_merge	0.17	1.00	0.83
p2_finance_fx	0.00	0.40	0.40
p3_ocr_invoice	0.00	1.00	1.00
p3_sql_recon	0.00	1.00	1.00
p4_eml_attachments	0.00	1.00	1.00
p4_xlsx_summary	0.00	0.00	0.00
Averages	0.05	0.84	0.79

Code Blame 22 lines (20 loc) · 899 Bytes

```
1  # Implement `run(pack_dir: str, prompt: str) -> dict | str`  
2  # - pack_dir: path to the located pack folder (so your agent can read files)  
3  # - prompt: same text as given to general LLM  
4  # Return JSON-compatible Python object, or a JSON string.  
5  #  
6  # TIP: You can route on task intent by reading files under `pack_dir`.  
7  # e.g., if 'answers.json' exists -> Pack1; 'answers_pack2.json' -> Pack2; etc.  
8  
9  import os, json  
10  
11 ✓ def run(pack_dir: str, prompt: str):  
12    # TODO: Wire your local tools. Below is a tiny heuristic demo:  
13    # This demo simply returns {}. Replace with real logic.  
14    # You can detect files under pack_dir and parse:  
15    # - CSVs in finance/, hr/, ops/  
16    # - .ics in calendar/  
17    # - PDFs in pdfs/ or pdf_scans/ (via OCR)  
18    # - PBM images in scans/ (Pack 3) -> OCR  
19    # - .eml and .mbox  
20    # - SQLite in product/ or sql/  
21    # - TAR/ZIP nested archives in archives/  
22    return {}
```

Next

More complex eval data (real-time data, larger files, huggingface repo)

New suite of data to challenge the memory conflicts of LLMs

- When local data doesn't match with the LLM's memory

More complex retrieval tasks (reasoning, and multi-step queries)

LLM researcher's toolkit

Tools:

- Training – Pytorch Lightning + Megatron (for LLM)
- Inference – vLLM, SGLang
- RL – veRL (Megatron + vLLM + Ray), SkyRL

Challenges/Opportunities:

- Inference merging into the training stage for RL
- System ⇒ more complexity, asynchronous, distributed ...
- Workload ⇒ multi-modal, long-sequences, sparse, streaming, dynamic


Goals of LLM Inference infrastructure

Speed – Achieve low latency (TTFT, TPOT)

Efficiency – Achieve higher throughput (tokens/sec) and lower cost (\$ / token)

Easy of Use – Support more models, Huggingface integration, API server ...

Introducing feedback loop



- Build a pipeline using the off-the-shelf frameworks
- Test out a few use cases and see the performance gain

Release 2.2 Roadmap – Inference feedbacks & User Interactions

We want to keep **data** on the edge as much as possible

EDA needs to understand the environment (w/ or w/o the help from the user)

When user starts to query the data, the agent retrieve the **information**

Ways of retrieval can be templated (but not-rule-based / hardcoded)

To-do-list:


Upper level application –? distill the **knowledge**

Update (fine-tune) the model for better services

Three Layers

1. Data
2. Information ← EDA
3. Knowledge


Demo (contributor: @ebowwa)




EDA for online-learning

Online Learning Demo

type in "label name" => take a photo on a new labeled data
sample => 512x512.png



Sleepetime compute



During sleep, the human brain sorts through different memories, consolidating important ones while discarding those that don't matter [https://www.wired.com/story/sleepetime-compute-chatbots-memory/?utm_source=chatgpt.com]

Memory-R1




Figure 2: Overview of the Memory-R1 framework. Stage 1 (blue) constructs and updates the memory bank via the RL-fine-tuned Memory Manager, which chooses operations {ADD, UPDATE, DELETE, NOOP} for each new dialogue turn. Stage 2 (green) answers user questions via the Answer Agent, which applies a Memory Distillation policy to reason over retrieved memories.

Idea

If you achieve the end result with IIM and record all the chats along the way, that means there is a path to there